

Using MobileNet for our Monkey Classifier

Loading the MobileNet Model

Freeze all layers except the top 4, as we'll only be training the top 4

```
In [8]: from keras.applications import MobileNet

# MobileNet was designed to work on 224 x 224 pixel input images sizes
img_rows, img_cols = 224, 224

# Re-loads the MobileNet model without the top or FC layers
MobileNet = MobileNet(weights = 'imagenet',
                        include_top = False,
                        input_shape = (img_rows, img_cols, 3))

# Here we freeze the last 4 layers
# Layers are set to trainable as True by default
for layer in MobileNet.layers:
    layer.trainable = False

# Let's print our layers
for (i, layer) in enumerate(MobileNet.layers):
    print(str(i) + " " + layer.__class__.__name__, layer.trainable)

0 InputLayer False
1 ZeroPadding2D False
2 Conv2D False
3 BatchNormalization False
4 ReLU False
5 DepthwiseConv2D False
6 BatchNormalization False
7 ReLU False
8 Conv2D False
9 BatchNormalization False
10 ReLU False
11 ZeroPadding2D False
12 DepthwiseConv2D False
13 BatchNormalization False
14 ReLU False
15 Conv2D False
16 BatchNormalization False
17 ReLU False
18 DepthwiseConv2D False
19 BatchNormalization False
20 ReLU False
21 Conv2D False
22 BatchNormalization False
23 ReLU False
24 ZeroPadding2D False
25 DepthwiseConv2D False
26 BatchNormalization False
27 ReLU False
28 Conv2D False
29 BatchNormalization False
30 ReLU False
31 DepthwiseConv2D False
32 BatchNormalization False
33 ReLU False
34 Conv2D False
35 BatchNormalization False
36 ReLU False
37 ZeroPadding2D False
38 DepthwiseConv2D False
39 BatchNormalization False
40 ReLU False
41 Conv2D False
42 BatchNormalization False
43 ReLU False
44 DepthwiseConv2D False
45 BatchNormalization False
46 ReLU False
47 Conv2D False
48 BatchNormalization False
49 ReLU False
50 DepthwiseConv2D False
51 BatchNormalization False
52 ReLU False
53 Conv2D False
54 BatchNormalization False
55 ReLU False
56 DepthwiseConv2D False
57 BatchNormalization False
58 ReLU False
59 Conv2D False
60 BatchNormalization False
61 ReLU False
62 DepthwiseConv2D False
63 BatchNormalization False
64 ReLU False
65 Conv2D False
66 BatchNormalization False
67 ReLU False
68 DepthwiseConv2D False
69 BatchNormalization False
70 ReLU False
71 Conv2D False
72 BatchNormalization False
73 ReLU False
74 ZeroPadding2D False
75 DepthwiseConv2D False
76 BatchNormalization False
77 ReLU False
78 Conv2D False
79 BatchNormalization False
80 ReLU False
81 DepthwiseConv2D False
82 BatchNormalization False
83 ReLU False
84 Conv2D False
85 BatchNormalization False
86 ReLU False
```

Let's make a function that returns our FC Head

```
In [9]: def lw(bottom_model, num_classes):
    """creates the top or head of the model that will be
    placed onto or the bottom layers"""

    top_model = bottom_model.output
    top_model = GlobalAveragePooling2D()(top_model)
    file_names = Dense(512, activation='relu')(top_model)
    top_model = Dense(256, activation='relu')(top_model)
    top_model = Dense(num_classes, activation='softmax')(top_model)
    return top_model
```

Let's add our FC Head back onto MobileNet

```
In [10]: from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, GlobalAveragePooling2D
from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D
from keras.layers.normalization import BatchNormalization
from keras.models import Model

# Set our class number to 3 (Young, Middle, Old)
num_classes = 3

FC_Head = lw(MobileNet, num_classes)

model = Model(inputs = MobileNet.input, outputs = FC_Head)

print(model.summary())

Model: "model_1"

Layer (type) Output Shape Param #
-----
input_1 (InputLayer) (None, 224, 224, 3) 0
conv1_pad (ZeroPadding2D) (None, 225, 225, 3) 0
conv1 (Conv2D) (None, 112, 112, 32) 864
conv1_bn (BatchNormalization) (None, 112, 112, 32) 128
conv1_relu (ReLU) (None, 112, 112, 32) 0
conv_dw_1 (DepthwiseConv2D) (None, 112, 112, 32) 288
conv_dw_1_bn (BatchNormalization) (None, 112, 112, 32) 128
conv_dw_1_relu (ReLU) (None, 112, 112, 32) 0
conv_pw_1 (Conv2D) (None, 112, 112, 64) 2048
conv_pw_1_bn (BatchNormalization) (None, 112, 112, 64) 256
conv_pw_1_relu (ReLU) (None, 112, 112, 64) 0
conv_pad_2 (ZeroPadding2D) (None, 113, 113, 64) 0
conv_dw_2 (DepthwiseConv2D) (None, 56, 56, 64) 576
conv_dw_2_bn (BatchNormalization) (None, 56, 56, 64) 256
conv_dw_2_relu (ReLU) (None, 56, 56, 64) 0
conv_pw_2 (Conv2D) (None, 56, 56, 128) 8192
conv_pw_2_bn (BatchNormalization) (None, 56, 56, 128) 512
conv_pw_2_relu (ReLU) (None, 56, 56, 128) 0
conv_dw_3 (DepthwiseConv2D) (None, 56, 56, 128) 1152
conv_dw_3_bn (BatchNormalization) (None, 56, 56, 128) 512
conv_dw_3_relu (ReLU) (None, 56, 56, 128) 0
conv_pw_3 (Conv2D) (None, 56, 56, 128) 16384
conv_pw_3_bn (BatchNormalization) (None, 56, 56, 128) 512
conv_pw_3_relu (ReLU) (None, 56, 56, 128) 0
conv_pad_4 (ZeroPadding2D) (None, 57, 57, 128) 0
conv_dw_4 (DepthwiseConv2D) (None, 28, 28, 128) 1152
conv_dw_4_bn (BatchNormalization) (None, 28, 28, 128) 512
conv_dw_4_relu (ReLU) (None, 28, 28, 128) 0
conv_pw_4 (Conv2D) (None, 28, 28, 256) 32768
conv_pw_4_bn (BatchNormalization) (None, 28, 28, 256) 1024
conv_pw_4_relu (ReLU) (None, 28, 28, 256) 0
conv_dw_5 (DepthwiseConv2D) (None, 28, 28, 256) 2304
conv_dw_5_bn (BatchNormalization) (None, 28, 28, 256) 1024
conv_dw_5_relu (ReLU) (None, 28, 28, 256) 0
conv_pw_5 (Conv2D) (None, 28, 28, 256) 65536
conv_pw_5_bn (BatchNormalization) (None, 28, 28, 256) 1024
conv_pw_5_relu (ReLU) (None, 28, 28, 256) 0
conv_pad_6 (ZeroPadding2D) (None, 29, 29, 256) 0
conv_dw_6 (DepthwiseConv2D) (None, 14, 14, 256) 2304
conv_dw_6_bn (BatchNormalization) (None, 14, 14, 256) 1024
conv_dw_6_relu (ReLU) (None, 14, 14, 256) 0
conv_pw_6 (Conv2D) (None, 14, 14, 512) 131072
conv_pw_6_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_pw_6_relu (ReLU) (None, 14, 14, 512) 0
conv_dw_7 (DepthwiseConv2D) (None, 14, 14, 512) 4608
conv_dw_7_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_dw_7_relu (ReLU) (None, 14, 14, 512) 0
conv_pw_7 (Conv2D) (None, 14, 14, 512) 262144
conv_pw_7_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_pw_7_relu (ReLU) (None, 14, 14, 512) 0
conv_dw_8 (DepthwiseConv2D) (None, 14, 14, 512) 4608
conv_dw_8_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_dw_8_relu (ReLU) (None, 14, 14, 512) 0
conv_pw_8 (Conv2D) (None, 14, 14, 512) 262144
conv_pw_8_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_pw_8_relu (ReLU) (None, 14, 14, 512) 0
conv_dw_9 (DepthwiseConv2D) (None, 14, 14, 512) 4608
conv_dw_9_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_dw_9_relu (ReLU) (None, 14, 14, 512) 0
conv_pw_9 (Conv2D) (None, 14, 14, 512) 262144
conv_pw_9_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_pw_9_relu (ReLU) (None, 14, 14, 512) 0
conv_dw_10 (DepthwiseConv2D) (None, 14, 14, 512) 4608
conv_dw_10_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_dw_10_relu (ReLU) (None, 14, 14, 512) 0
conv_pw_10 (Conv2D) (None, 14, 14, 512) 262144
conv_pw_10_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_pw_10_relu (ReLU) (None, 14, 14, 512) 0
conv_dw_11 (DepthwiseConv2D) (None, 14, 14, 512) 4608
conv_dw_11_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_dw_11_relu (ReLU) (None, 14, 14, 512) 0
conv_pw_11 (Conv2D) (None, 14, 14, 512) 262144
conv_pw_11_bn (BatchNormalization) (None, 14, 14, 512) 2048
conv_pw_11_relu (ReLU) (None, 14, 14, 512) 0
conv_pad_12 (ZeroPadding2D) (None, 15, 15, 512) 0
conv_dw_12 (DepthwiseConv2D) (None, 7, 7, 512) 4608
conv_dw_12_bn (BatchNormalization) (None, 7, 7, 512) 2048
conv_dw_12_relu (ReLU) (None, 7, 7, 512) 0
conv_pw_12 (Conv2D) (None, 7, 7, 1024) 524288
conv_pw_12_bn (BatchNormalization) (None, 7, 7, 1024) 4096
conv_pw_12_relu (ReLU) (None, 7, 7, 1024) 0
conv_dw_13 (DepthwiseConv2D) (None, 7, 7, 1024) 9216
conv_dw_13_bn (BatchNormalization) (None, 7, 7, 1024) 4096
conv_dw_13_relu (ReLU) (None, 7, 7, 1024) 0
conv_pw_13 (Conv2D) (None, 7, 7, 1024) 1048576
conv_pw_13_bn (BatchNormalization) (None, 7, 7, 1024) 4096
conv_pw_13_relu (ReLU) (None, 7, 7, 1024) 0
global_average_pooling2d_1 ( (None, 1024) 0
dense_1 (Dense) (None, 512) 524800
dense_2 (Dense) (None, 256) 131328
dense_3 (Dense) (None, 3) 771
Total params: 3,885,763
Trainable params: 655,899
Non-trainable params: 3,228,864
None
```

Loading our Monkey Breed Dataset

```
In [11]: from keras.preprocessing.image import ImageDataGenerator

train_data_dir = 'face_recognition/face_recognition/train/'
validation_data_dir = 'face_recognition/face_recognition/validation/'

# Let's use some data augmentation
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=45,
    width_shift_range=0.3,
    height_shift_range=0.3,
    horizontal_flip=True,
    fill_mode='nearest')

validation_datagen = ImageDataGenerator(rescale=1./255)

# set our batch size (typically on most mid tier systems we'll use 16-32)
batch_size = 32

train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(img_rows, img_cols),
    batch_size=batch_size,
    class_mode='categorical')

validation_generator = validation_datagen.flow_from_directory(
    validation_data_dir,
    target_size=(img_rows, img_cols),
    batch_size=batch_size,
    class_mode='categorical')

Found 30 images belonging to 3 classes.
Found 20 images belonging to 3 classes.
```

Training our Model

- Note we're using checkpointing and early stopping

```
In [12]: from keras.optimizers import RMSprop
from keras.callbacks import ModelCheckpoint, EarlyStopping

checkpoint = ModelCheckpoint("face_recognition_mobileNet.h5",
                             monitor="val_loss",
                             mode="min",
                             save_best_only = True,
                             verbose=1)

earlystop = EarlyStopping(monitor = 'val_loss',
                           min_delta = 0,
                           patience = 3,
                           verbose = 1,
                           restore_best_weights = True)

# We put our call backs into a callback list
callbacks = [earlystop, checkpoint]

# We use a very small learning rate
model.compile(loss = 'categorical_crossentropy',
              optimizer = RMSprop(lr = 0.001),
              metrics = ['accuracy'])

# Enter the number of training and validation samples here
nb_train_samples = 1097
nb_validation_samples = 272

# We only train 5 EPOCHS
epochs = 3
batch_size = 16

history = model.fit_generator(
    train_generator,
    steps_per_epoch = nb_train_samples // batch_size,
    epochs = epochs,
    callbacks = callbacks,
    validation_data = validation_generator,
    validation_steps = nb_validation_samples // batch_size)

Epoch 1/3
68/68 [=====] - 374s 6s/step - loss: 0.2156 - accuracy: 0.9657 - val
_loss: 1.8606 - val_accuracy: 0.6500

Epoch 00001: val_loss improved from inf to 1.86064, saving model to face_recognition_mobileNe
t.h5
31/68 [=====] - ETA: 2:37 - loss: 2.2637e-05 - accuracy: 1.0000

KeyboardInterrupt (pid=1440d4573dac) in <module>: Traceback (most recent call last)
--> 38 validation_data = validation_generator,
--> 39 validation_steps = nb_validation_samples // batch_size)

E:\python\python\lib\site-packages\keras\legacy\interfaces.py in wrapper(*args, **kwargs)
    89 warnings.warn("Update your '" + object_name + "' call to the '" +
    90 keras 2 API: " + signature, stacklevel=2)
--> 91 return func(*args, **kwargs)
    92 wrapper._original_function = func
    93 return wrapper

E:\python\python\lib\site-packages\keras\engine\training.py in fit_generator(self, generator,
steps_per_epoch, epochs, verbose, callbacks, validation_data, validation_steps, val
idation_freq, class_weight, max_queue_size, workers, use_multiprocessing, shuffle, initial_e
poch)
    1730 use_multiprocessing=use_multiprocessing,
--> 1732 shuffle=shuffle,
    1733 initial_epoch=initial_epoch)
    1734 @interfaces.legacy_generator_methods_support

E:\python\python\lib\site-packages\keras\engine\training_generator.py in fit_generator(model,
generator, steps_per_epoch, epochs, verbose, callbacks, validation_data, validation_steps, va
lidation_freq, class_weight, max_queue_size, workers, use_multiprocessing, shuffle, initial_e
poch)
    218 sample_weight=sample_weight,
    219 class_weight=class_weight,
--> 220 reset_metrics=False)
    221
    222 outs = to_list(outs)

E:\python\python\lib\site-packages\keras\engine\training.py in train_on_batch(self, x, y, sam
ple_weight, class_weight, reset_metrics)
    1512 ins = x + y + sample_weights
    1513 self._make_train_function()
--> 1514 outputs = self.train_function(ins)
    1515
    1516 if reset_metrics:

E:\python\python\lib\site-packages\tensorflow_core\python\keras\backend.py in __call__(self,
inputs)
    3725 value = math_ops.cast(value, tensor.dtype)
--> 3727 outputs = self._graph_fn(*converted_inputs)
    3728
    3729 # EagerTensor.numpy() will often make a copy to ensure memory safety.

E:\python\python\lib\site-packages\tensorflow_core\python\eager\function.py in __call__(self,
*args, **kwargs)
    1549 TypeError: For invalid positional/keyword argument combinations.
--> 1551 return self._call_impl(args, kwargs)
    1552
    1553 def _call_impl(self, args, kwargs, cancellation_manager=None):

E:\python\python\lib\site-packages\tensorflow_core\python\eager\function.py in _call_impl(sel
f, args, kwargs, cancellation_manager)
    1589 raise TypeError("keyword arguments {} unknown. Expected {}.".format(
    1590 list(kwargs.keys()), list(self.arg_keywords)))
--> 1591 return self._call_flat(args, self.captured_inputs, cancellation_manager)
    1592
    1593 def _filtered_call(self, args, kwargs):

E:\python\python\lib\site-packages\tensorflow_core\python\eager\function.py in _call_flat(sel
f, args, captured_inputs, cancellation_manager)
    1690 # No tape is watching; skip to running the function.
    1691 return self._build_call_outputs(self._inference_function.call(
--> 1692 ctx, args, cancellation_manager=cancellation_manager))
    1693 forward_backward = self._select_forward_and_backward_functions(
    1694 args,

E:\python\python\lib\site-packages\tensorflow_core\python\eager\function.py in call(self, ct
x, args, cancellation_manager)
    543 inputs=args,
    544 attr=("executor_type", executor_type, "config_proto", config),
--> 545 ctx=ctx)
    546 else:
    547 outputs = execute.execute_with_cancellation(

E:\python\python\lib\site-packages\tensorflow_core\python\eager\execute.py in quick_execute(o
p_name, num_outputs, inputs, attrs, ctx, name)
    59 tensors = pywrap_tensorflow.TFE_Py_Execute(ctx._handle, device_name,
--> 60 op_name, inputs, attrs,
    61 num_outputs)
    62 except core._NotOkStatusException as e:
    63 if name is not None:

KeyboardInterrupt:

Class - Vikas
```

Loading our classifier

```
In [13]: from keras.models import load_model

classifier = load_model('face_recognition_mobileNet.h5')
```

Testing our classifier on some test images

```
In [ ]: import os
import cv2
import numpy as np
from os import listdir
from os.path import isfile, join

face_recognition_dict = {"[0]": "Abhinay",
                        "[1]": "Abhishek",
                        "[2]": "Vikas"}

face_recognition_dict_n = {"Abhinay": "Abhinay",
                          "Abhishek": "Abhishek",
                          "Vikas": "Vikas"}

def draw_test(name, pred, im):
    face = face_recognition_dict[str(pred)]
    BLACK = [0,0,0]
    expanded_image = cv2.cvtColor(makeBorder(im, 80, 0, 0, 180, cv2.BORDER_CONSTANT, value=BLACK)
    cv2.putText(expanded_image, face, (20, 60), cv2.FONT_HERSHEY_SIMPLEX, 1, (0,0,255), 2)
    cv2.imshow(name, expanded_image)

def getRandomImage(path):
    """function loads a random images from a random folder in our test path"""
    folders = list(filter(lambda x: os.path.isdir(os.path.join(path, x)), os.listdir(path)))
    random_dir = random.choice(folders)
    path_class = folders[random_dir]
    print("Class - " + face_recognition_dict_n[str(path_class)])
    file_path = " " + path_class
    file_names = [f for f in listdir(file_path) if isfile(join(file_path, f))]
    random_file_index = np.random.randint(0, len(file_names))
    image_name = file_names[random_file_index]
    return cv2.imread(file_path+"/"+image_name)

for i in range(0,3):
    input_im = getRandomImage("face_recognition/face_recognition/validation/")
    input_original = input_im.copy()
    input_original = cv2.resize(input_original, None, fx=0.5, fy=0.5, interpolation = cv2.INTER_LINEAR)

    # Get Prediction
    res = np.argmax(classifier.predict(input_im, 1, verbose = 0), axis=1)

    # Show image with predicted class
    draw_test("Prediction", res, input_original)
    cv2.waitKey(0)

cv2.destroyAllWindows()
```

Class - Vikas